

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

In [2]:

```
df=pd.read_csv(r"C:\Users\raja\Downloads\fiat500_VehicleSelection_Dataset.csv")
df
```

Out[2]:

| | ID | model | engine_power | age_in_days | km | previous_owners | lat | lon | price |
|------|------|--------|--------------|-------------|--------|-----------------|-----------|-----------|-------|
| 0 | 1 | lounge | 51 | 882 | 25000 | 1 | 44.907242 | 8.611560 | 8900 |
| 1 | 2 | pop | 51 | 1186 | 32500 | 1 | 45.666359 | 12.241890 | 8800 |
| 2 | 3 | sport | 74 | 4658 | 142228 | 1 | 45.503300 | 11.417840 | 4200 |
| 3 | 4 | lounge | 51 | 2739 | 160000 | 1 | 40.633171 | 17.634609 | 6000 |
| 4 | 5 | pop | 73 | 3074 | 106880 | 1 | 41.903221 | 12.495650 | 5700 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1533 | 1534 | sport | 51 | 3712 | 115280 | 1 | 45.069679 | 7.704920 | 5200 |
| 1534 | 1535 | lounge | 74 | 3835 | 112000 | 1 | 45.845692 | 8.666870 | 4600 |
| 1535 | 1536 | pop | 51 | 2223 | 60457 | 1 | 45.481541 | 9.413480 | 7500 |
| 1536 | 1537 | lounge | 51 | 2557 | 80750 | 1 | 45.000702 | 7.682270 | 5990 |
| 1537 | 1538 | pop | 51 | 1766 | 54276 | 1 | 40.323410 | 17.568270 | 7900 |

1538 rows × 9 columns

In [3]:

df.columns

Out[3]:

```
Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners',
      'lat', 'lon', 'price'],
      dtype='object')
```

In [4]:

```
x=df[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners', 'lat', 'lon']]
y=df['price']
```

In [5]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print(regr.score(x_test,y_test))
```

0.8442339248908175

In [6]:

```
lm=LinearRegression()
lm.fit(x_train,y_train)
```

Out[6]:

LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [7]:

```
coeff_df=pd.DataFrame(lm.coef_,x.columns,columns=['coefficient'])
coeff_df
```

Out[7]:

| | coefficient |
|-----------------|-------------|
| ID | -0.055110 |
| engine_power | 11.077477 |
| age_in_days | -0.900816 |
| km | -0.017422 |
| previous_owners | 7.287449 |
| lat | 37.713807 |
| lon | 5.681090 |

In [8]:

```
features = ['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',
            'lat', 'lon']
```

In [9]:

```
target=['price']
```

In [10]:

```
from sklearn.linear_model import Ridge,RidgeCV,Lasso
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
```

In [11]:

```
ridgeReg = Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
train_score_ridge = ridgeReg.score(x_train,y_train)
test_score_ridge = ridgeReg.score(x_test,y_test)
print('\nRidge model\n')
print('Train score for ridge model is {}'.format(train_score_ridge))
print('Test score for ridge model is {}'.format(test_score_ridge))
```

Ridge model


```
Train score for ridge model is 0.8421192197724451
Test score for ridge model is 0.8442304722432865
```

In [14]:

```

t.figure(figsize=(10,10))
t.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=7,color='Red',label=r'ridge;$\alpha=10$',zorder=7)
t.plot(features,lm.coef_,alpha=0.4,linestyle='none',color='green',marker='d',markersize=6,label='LinearRegression')
t.xticks(rotation=90)
t.legend()
t.show()

```



In [15]:

```

lassoReg=Lasso(alpha=10)
lassoReg.fit(x_train,y_train)
train_score_lasso=lassoReg.score(x_train,y_train)
test_score_lasso=lassoReg.score(x_test,y_test)
print('\nLasso Model\n')
print('Train score for lasso model is {}'.format(train_score_lasso))
print('Test score for lasso model is {}'.format(test_score_lasso))


```

Lasso Model

Train score for lasso model is 0.8420908335130204
 Test score for lasso model is 0.8441608351601755

In [16]:

```
t.figure(figsize=(10,10))
t.plot(features,ridgeReg.coef_,alpha=0.5,linestyle='none',marker='*',markersize=6,color='red',label=r'lasso;$\alpha=10$',zorder=7)
t.plot(features,lm.coef_,alpha=0.6,linestyle='none',marker='d',markersize=7,color='k',label='LinearRegression')
t.xticks(rotation=90)
t.legend()
t.show()
```



In [17]:

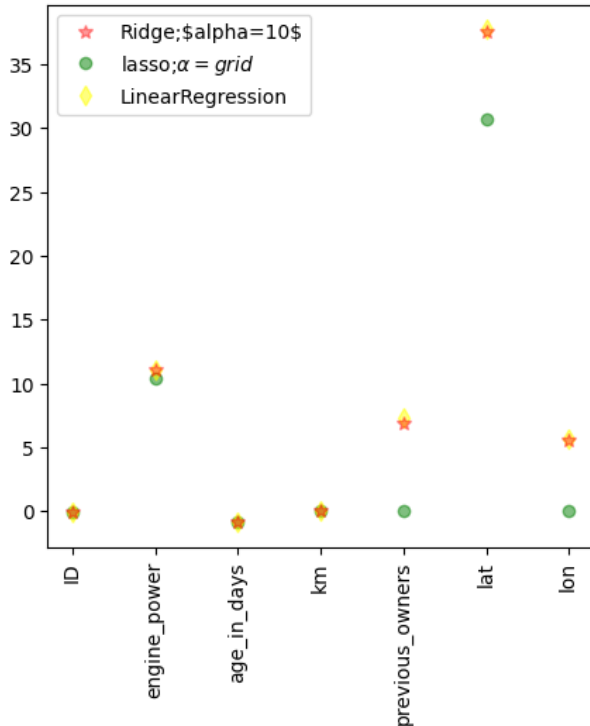
comaprison between Ridge,Lasso and ridgeCV

```

t.figure(figsize=(5,5))
t.plot(features,ridgeReg.coef_,alpha=0.4,linestyle='none',marker='*',markersize=7,color='red',label=r'Ridge;\alpha=10$',zorder=7)
t.plot(features,lassoReg.coef_,alpha=0.5,linestyle='none',marker='o',markersize=6,color='green',label=r'lasso;\alpha=grid$')
t.plot(features,lm.coef_,alpha=0.5,linestyle='none',marker='d',markersize=7,color='yellow',label='LinearRegression')
t.xticks(rotation=90)
t.title('Comparison between Rudge,Lasso and RidgeCV')
t.legend()
t.show()

```

Comparison between Rudge,Lasso and RidgeCV



In [18]:

Linear CV model using Ridge

```

from sklearn.linear_model import RidgeCV
ridge_CV=RidgeCV(alphas=[0.1,0.4,1.1]).fit(x_train,y_train)
print('The Train score for ridge model is {}'.format(ridge_CV.score(x_train,y_train)))
print('The Test score for ridge model is {}'.format(ridge_CV.score(x_test,y_test)))

```

The Train score for ridge model is 0.8421187226355313

The Test score for ridge model is 0.8441837688019272

In [19]:

Linear CV model using Lasso

```

from sklearn.linear_model import LassoCV
lasso_CV=LassoCV(alphas=[1,10,20]).fit(x_train,y_train)
print("The train score for lasso model is {}".format(lasso_CV.score(x_train,y_train)))
print("The test score for lasso model is {}".format(lasso_CV.score(x_test,y_test)))

```

The train score for lasso model is 0.8420675579097567

The test score for lasso model is 0.8441521459676745

In [22]:

Elstic Net

```

from sklearn.linear_model import ElasticNet
regr = ElasticNet()
regr.fit(x,y)
#print(regr.coef_)
print(regr.intercept_)
y_pred_elastic = regr.predict(x_train)
mean_squared_error = np.mean((y_pred_elastic-y_train)**2)
print('Mean squared error on test set',mean_squared_error)

```

9248.105972029809

Mean squared error on test set 593118.858587302

